

Three weeks ago, Copilot started billing by the token.

On **June 1, 2026** Copilot moved to usage-based **AI Credits**. Every chat, edit, and agent request is now metered on input + output + cached tokens at the model's own rate. **Token efficiency is no longer a nicety — it's the invoice.** Here is the team playbook to cut that bill without losing quality.

Jun 1 Switched to usage-based AI Credits	tokens = \$ Input + output + cached, per model rate
10–20× Cost spread, frugal vs top-tier model	\$0 Inline completions on paid plans

THE 6 LEVERS, RANKED BY IMPACT

Where the money actually goes

Pull these top to bottom. The first two — model choice and context scope — move the bill more than the rest combined. Everything here is Copilot-native, not generic advice.

Match the model to the task

1

The dominant lever. A frugal model can cost ~10–20× less per token for the same routine work. Prefer **auto-selection** — it routes by task and discounts the rate. Never run a reasoning model on a rename.

IMPACT



Scope the context

2

Input tokens are billed every request. Attach specific files with `#file` / `#selection` instead of `@workspace` for narrow tasks — and **start a fresh chat per task**, since long threads re-bill stale context on every turn.

IMPACT



Pick the cheapest sufficient mode

3

Cost climbs `completions → ask → edit → agent`. Completions are free; agent loops re-send context and tools each turn and quietly become the largest single line item. Reserve agent for genuinely multi-step work.

IMPACT



Cut output verbosity

4

Output tokens bill too. A standing “be concise, prefer diffs over full-file rewrites” instruction trims every future response. Set it once in `copilot-instructions.md`.

IMPACT



Reuse, don't re-type

5

Stop paying to re-establish context. `*.instructions.md` (with `applyTo` globs) and `*.prompt.md` load standing rules and recurring prompts automatically instead of re-sending them by hand.

IMPACT



Make code review deliberate

6

Automated Copilot code review spends credits *and* GitHub Actions minutes per run. Scope it to meaningful PRs and paths rather than firing on every push.

IMPACT



LEVER 1, IN PRACTICE

The model-selection matrix

The single most useful habit to teach. Default down the list; reach up only when the task earns it. Cost shown is relative per-token rate, not a quota.

TASK	USE	RELATIVE COST
Autocomplete, boilerplate, obvious edits	● Inline completions No chat request at all	FREE
Simple Q&A, rename, small fix, docstring	● Frugal / fast GPT-5 mini · Gemini Flash · Claude Haiku	\$
Day-to-day chat, scoped edits, reviews	● Standard Mid Sonnet · GPT-5 class	\$\$

TASK	USE	RELATIVE COST
Hard architecture, deep debugging, big refactors	● High / reasoning Claude Opus · GPT-5.5 — only when earned	\$\$\$\$

LEVERS 4 & 5 · SHIP THESE MONDAY

Three files that pay for themselves

Drop these in any repo. They load automatically, so the whole team stops paying to re-establish conventions and context in every thread.

`.github/copilot-instructions.md`

Repo-wide standing rules

"Lead with the answer. Prefer diffs over full-file rewrites. Be terse for experts." Applies to every chat, edit, and agent request.

CUTS OUTPUT TOKENS

`.github/instructions/*.instructions.md`

Path-specific rules

Frontmatter `applyTo: "**/*.test.*"` targets a glob. Matching files **stack** (union, no override) — keep each one small.

SCOPES CONTEXT

`.github/prompts/*.prompt.md`

Reusable prompts

A saved review checklist or scaffold. Invoke it instead of re-typing — and re-paying for — a long prompt every time.

REUSE, DON'T RE-TYPE

Don't save credits by getting it wrong

Cheaper-but-wrong is the most expensive outcome — it costs a second request, a bad merge, or an outage. Compression never touches these:

exact paths & commands

error strings

identifiers & versions

security warnings

material risks

validation status

the actual answer

Priority when cost conflicts with quality: **safety → correctness → user intent → validation → credit savings.**

Prove it, don't assume it

Track AI Credit usage in GitHub billing before and after — org cost centers and budgets land on Business / Enterprise. A-B one real task (verbose baseline vs. this config, same model) and score quality 1–5 separately; adopt only if quality holds.

Plan credits (monthly, as of Jun 2026): Pro \$10 · Pro+ \$39 · Business \$19/user · Enterprise \$39/user. Free = completions only.

Source: AgentDefaults · [skills/copilot-token-efficiency.md](#) · [examples/copilot-token-efficiency.md](#)

Verify before you quote. GitHub changes model rates, multipliers, and plan allowances often, and the move off legacy premium-request billing is recent. Confirm current numbers at [docs.github.com → Copilot billing → models & pricing](#) before presenting figures as fact.